

ASIM AWAD HUSSEIN OSMAN - Needed help from Tim Hitge at section 2 part regarding the debug of the overtraining plot code, I was not using the range parameter in the histogram. and I was not sure about the separation of the signal/background that should be used to plot the histogram bars. other than that, I worked alone.

✓ Part 1. Decision Stump by hand

✓ 1.1 Get and load the data

Get the train dataset and load the relevant columns in a dataframe.

```
# imports
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
# mount Google Drive
from google.colab import drive
drive.mount('/content/gdrive')
```

↗ Mounted at /content/gdrive

```
# prepare paths
train_path = "/content/gdrive/MyDrive/data_atlas_higgs_4lep/ATLAS_higgs_train.csv"
test_path  = "/content/gdrive/MyDrive/data_atlas_higgs_4lep/ATLAS_higgs_test.csv"
# import the data to dataframes
train_df = pd.read_csv(train_path)
test_df  = pd.read_csv(test_path)
```

```
# Explore the data
train_df.describe(),train_df.head(10),train_df.columns
# test_df.describe(),test_df.head(10),test_df.columns
```

↗ min 7.030563 -13.000000 -2.699395 -3.141244 7.000066

```

5  00.214414      11  0.307432 ...  40.239031      -13  0.190140  0.223073
6  71.828297     -13  1.223559 ...  36.534945      13  0.135275  1.105578
7  54.883793      13  0.399076 ...  48.801734     -13  0.013284  0.096804
8  89.726109      11 -0.113298 ...  71.146961     -13 -0.490779 -0.016436
9  49.593051      11 -0.925869 ...  40.077438     -13 -1.581438  2.941436

```

```

      l3pt  l4pdgid      l4eta      l4phi      l4pt  sample
0  28.700672      11  0.542029 -0.328558  17.054512      1
1  30.645279      13  1.380638  0.300915   8.034711     -1
2  17.942230     -11  0.672241 -1.410558  14.799019     -1
3  12.816265      13 -2.283192 -2.326921  11.448904      1
4  22.981393      13  0.099710  1.758643  12.038583      1
5  19.296023      13  0.979412  1.690349  14.323106      1
6  22.613949      13  0.420407  2.463376  21.033371      1
7  18.417072     -13 -1.192064  1.455519  12.260007      1
8  13.056511      13 -1.359360 -2.184997   8.228871      1
9  13.862145      13 -1.713953 -2.456506   8.767364     -1

[10 rows x 21 columns],
Index(['Z1mass', 'Z2mass', 'detajj', 'massjj', 'l1pdgid', 'l1eta', 'l1phi',
      'l1pt', 'l2pdgid', 'l2eta', 'l2phi', 'l2pt', 'l3pdgid', 'l3eta',
      'l3phi', 'l3pt', 'l4pdgid', 'l4eta', 'l4phi', 'l4pt', 'sample'],
      dtype='object'))

```

```
# Data processing and setting global variables
```

```
# GLOBAL VARIABLES
```

```

XNAME = 'detajj'; XLABEL = r'$|\Delta\eta_{jj}|$'
YNAME = 'massjj'; YLABEL = r'$m_{jj}$ (GeV)'
TARGET = 'sample'
inputs= [XNAME, YNAME] ;

```

```

XBINS = 5 ; XMIN = 0 ; XMAX = 5      ; XSTEP = 0.1
YBINS = 5 ; YMIN = 0 ; YMAX = 1000 ; YSTEP = 10
FONTSIZE = 16

```

```

# preparing the range at which the cut of features will take place (making steps along the range of feature) will be used at th
detajj_range = np.arange(XMIN,XMAX,XSTEP)
massjj_range = np.arange(YMIN,YMAX,YSTEP)

```

```
# Creating reduced datasets with detajj & massjj only
```

```

train_df = train_df[['detajj', 'massjj','sample']]
test_df = test_df[['detajj', 'massjj','sample']]

```

```
# splitting data to X,y fashion
```

```

X_train = train_df[inputs] ; y_train = train_df[['sample']]
X_test = test_df[inputs] ; y_test = test_df[['sample']]

```

```
# splittin the signal from the background in both of the train and test data
```

```
# For Training
```

```

Signal_train = train_df[train_df[TARGET] == 1]
Background_train = train_df[train_df[TARGET] == -1]

```

```
# For test
```

```

Signal_test = test_df[test_df['sample'] == 1][inputs]
Background_test = test_df[test_df['sample'] == -1][inputs]

```

```
type(y_train)
```



```

pandas.core.frame.DataFrame
def __init__(data=None, index: Axes | None=None, columns: Axes | None=None, dtype: Dtype |
None=None, copy: bool | None=None) -> None

>>> d = {'col1': [0, 1, 2, 3], 'col2': pd.Series([2, 3], index=[2, 3])}
>>> pd.DataFrame(data=d, index=[0, 1, 2, 3])
   col1  col2
0      0   NaN
1      1   NaN
2      2   2.0

```

✓ 1.2 Compute the Gini index

Write a function computing the Gini index value. Make your code as general as possible. Add in the next cells a series of tests to make sure your function returns the correct answers.

Bonus: secure your code to prevent a division by zero.

```
# function to calculate the Gini index, it takes the dataframe, returns the index
def Gini_index(dataFrame ,target_class):

    total_samples = dataFrame.shape[0]

    #if there are no samples - prevent division by zero
    if total_samples == 0:
        return 0

    # get the unique classes available at the target ('sample') column
    classes = dataFrame[target_class].unique()
    classes_dist = {}
    # count the number of occurrences of each class
    for cls in classes:
        classes_dist[cls] = dataFrame[dataFrame[target_class] == cls].shape[0]
    # print(classes_dist)

    #calculate the gini index
    gini = 1
    for i in classes:
        gini -= (classes_dist[i]/total_samples)**2
    return gini
```

Gini Index Tests

```
# avoiding division by zero
df_zero = pd.DataFrame(columns= ['sample'])
print(Gini_index(df_zero,'sample'))

# the entire data set should be 50-50 number of samples, so a result would be 0.5
print(Gini_index(train_df,'sample'))

# another random data to calculate gini by hand (0.48)
dft = pd.DataFrame({'detajj': [10, 12, 43, 54, 665],
                    'sample': [0, 0, 0, 1, 1]
})
(Gini_index(dft, 'sample'))
```

```
0
0.5
0.48
```

1.3 Calculate the cost

Write a function computing the cost function in the CART algorithm. Again test it with a dummy dataset/example and comment.

Bonus: secure your code to prevent a division by zero.

```
def calculate_split_cost(dataFrame,feature, threshold , target_class):

# split the data frame according the feature and threshold
left_df = dataFrame[dataFrame[feature] >= threshold]
df_right = dataFrame[dataFrame[feature] < threshold]

# calculate the numbers of samples in each class and the total number of samples
N_left = left_df.shape[0]
# print(N_left)
N_right= df_right.shape[0]
# print(N_right)

N_Node = N_left + N_right

# check division by zero
if N_Node == 0 :
    return 0
# calculate the cost multiply gini from the left by the number of samples on the left multiplied by the total samples for the right
```

```

# calculate the cost multiply gini from the left by the number of samples on the left multiplied by the total - same for the r
cost = ( (N_left * Gini_index(left_df,target_class)) + (N_right * Gini_index(df_right,target_class)) )/N_Node

return cost

```

Example testing the function

```

# usnig the entire traingin set
print(calculate_split_cost(train_df,feature= XNAME, threshold = 990, target_class = 'sample'))
# usin gthe previous example from gini index
dft = pd.DataFrame({'detajj': [10, 12, 43, 54, 665],
                    'sample': [0, 1, 0, 1, 1]
})
print(calculate_split_cost(dft,feature= "detajj", threshold = 50, target_class = 'sample'))
# # this should be a 2.6 cost

0.5
0.2666666666666667

```

✓ 1.4 Main function: code a Decision Stump

Write the main function decision_stumper that will call the functions defined above.

Call decision_stumper on each input feature

Conclude on the final cut for your decision stump (explain your reasoning)

Optional: if you want, you can enter a list of features and do the two steps above in the same function. Comment your code appropriately.

```

# this function takes a feature and a dataframe and gives the best position to cut along the range of this feature, the split v

def Find_best_feature_cut(dataFrame,feature, thresholds ,target_class):
    # prepare a list to store the (cost,threshold) pairs
    costs = []
    cut = 0 # the cut starting position
    # loop over all possible cuts on the feature at hand ( we will go on a step size depending on the feature itslef)
    for point in thresholds:
        split_cost = calculate_split_cost(dataFrame,feature, point, target_class)
        # append the cost associated with this threshold to the list of costs
        costs.append((split_cost,point))
    #Finally return the threshold that corresponds to the least cost along with the cut position (threshold)
    return min(costs)

Find_best_feature_cut(dataFrame=train_df,feature ='detajj', thresholds = detajj_range ,target_class ='sample')

(0.4659133723626246, 2.6)

# features_thresholds = {'detajj' : detajj_range , 'massjj':massjj_range}

# features_thresholds.items()

# Setting Some Variables
# dictionary containing the features with their ranges (problem specific to this assignment)
features_thresholds = {'detajj' : detajj_range , 'massjj':massjj_range}

def DECISION_STUMPER(dataFrame , features_thresholds , target_column):

    # list to store the costs across different features
    overall_costs = []

    # loop over every feature getting the lowest cost cut accross each, then calculating the minimum of them as the primary cut fc
    for feature,range in features_thresholds.items():
        cost , threshold = Find_best_feature_cut(dataFrame,feature, range, target_column)
        # add this cost to the overal cost list
        overall_costs.append((cost,threshold,feature))
        # Printout to show the best cuts accross each feature
        # print(f"feature : {feature}, cost : {cost} , threshold : {threshold}")

    # Move one to perform splitting the data frame to tow dataframes based on the minimum cost calculated
    final_cut_informations = min(overall_costs)

```

```

cost = final_cut_informations[0] # the minimum cost accross all features
cut_threshold = final_cut_informations[1] # the most discriminative threshold
cut_feature = final_cut_informations[2] #the most discriminative feature (one chosen)

# now perform the split
left_df = dataframe[dataframe[cut_feature] >= cut_threshold]
df_right = dataframe[dataframe[cut_feature] < cut_threshold]
results = {'Best Feature' : cut_feature , 'best threshold ' : cut_threshold , 'best cost' : cost }

return results

results = DECISION_STUMPER(train_df , features_thresholds, target_column='sample')
print(results)
optimal_threshold = results['best threshold ']

➦ {'Best Feature': 'massjj', 'best threshold ': 220, 'best cost': 0.45639043953231406}

```

✓ 1.5 Plot the cut

Use the plot_scatter function from the second tutorial and modify it to draw the line corresponding to the optimized threshold from the decision stump. You can use Matplotlib's axhline or axvline method for drawing a horizontal or vertical line respectively. Try to be as general as possible in the input arguments.

```

def plot_scatter(sig, bkg,best_threshold,
                 xname=XNAME, xlabel=XLABEL, xmin=XMIN, xmax=XMAX, xstep=1,
                 yname=YNAME, ylabel=YLABEL, ymin=YMIN, ymax=YMAX, ystep=200,
                 fgsz=(6, 6), ftsz=16, alpha=0.3, title="Scatter plot"):

    fig, ax = plt.subplots(figsize=fgsz)

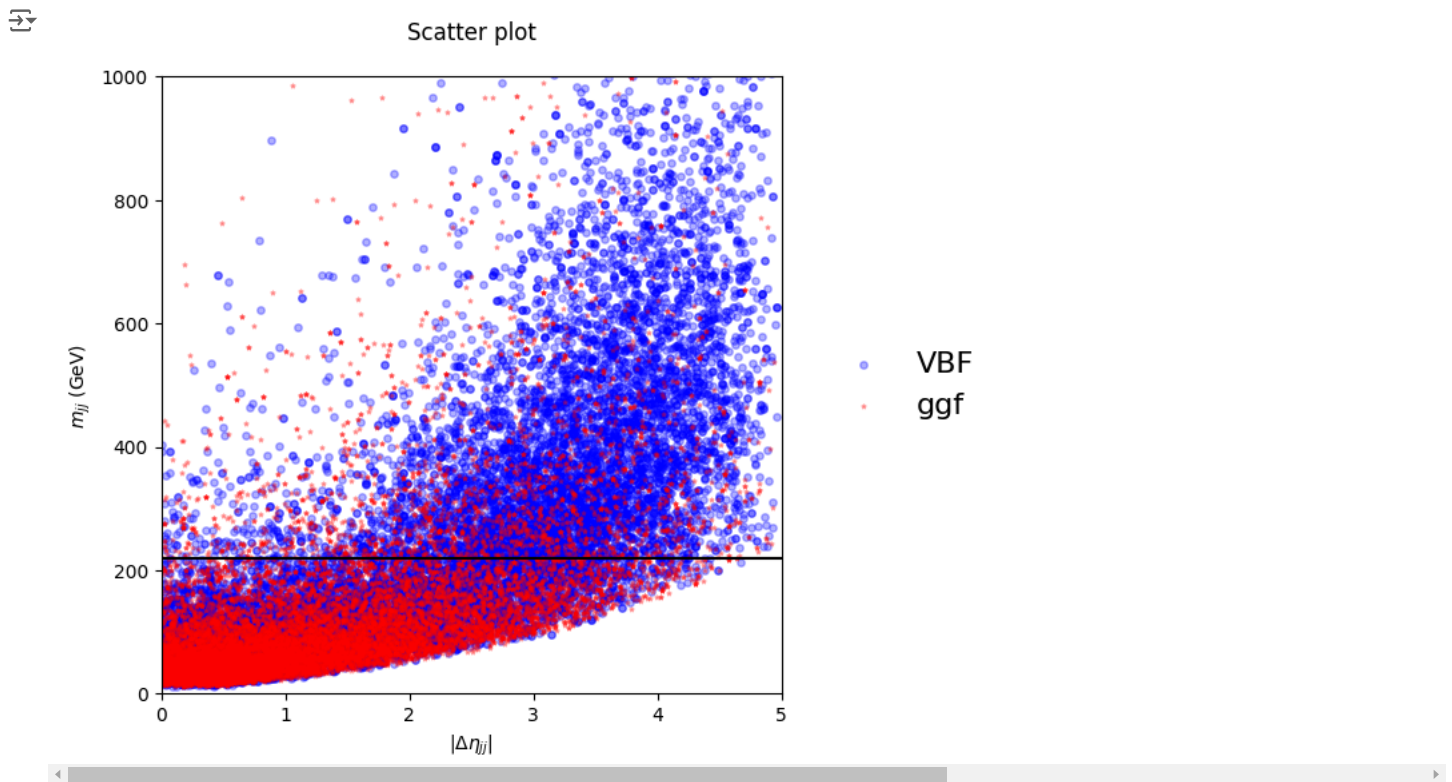
    # Annotate x-axis
    ax.set_xlim(xmin, xmax)
    ax.set_xlabel(xlabel)
    ax.set_xticks(np.arange(xmin, xmax+xstep, xstep))

    # Annotate y-axis
    ax.set_ylim(ymin, ymax)
    ax.set_ylabel(ylabel)
    ax.set_yticks(np.arange(ymin, ymax+ystep, ystep))

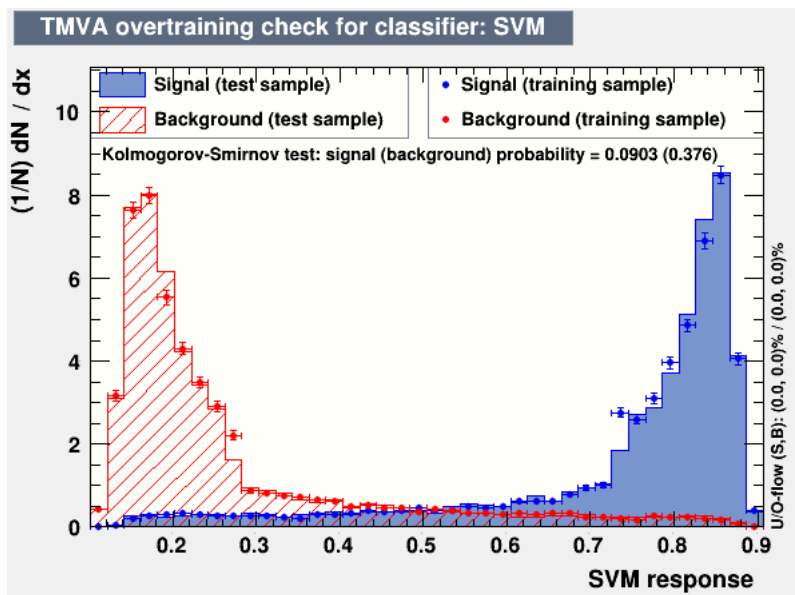
    # Scatter signal and background:
    ax.scatter(sig[xname], sig[yname], marker='o', s=15, c='b', alpha=alpha, label='VBF')
    ax.scatter(bkg[xname], bkg[yname], marker='*', s= 5, c='red', alpha=alpha, label='ggf')
    plt.axhline(y = best_threshold, c = 'black')
    # Legend and plot:
    ax.legend(fontsize=ftsz, bbox_to_anchor=(1.04, 0.5), loc="center left", frameon=False)
    ax.set_title(title, pad=20)
    plt.show()

plot_scatter(Signal_train, Background_train,optimal_threshold,
            xname=XNAME, xlabel=XLABEL, xmin=XMIN, xmax=XMAX, xstep=1,
            yname=YNAME, ylabel=YLABEL, ymin=YMIN, ymax=YMAX, ystep=200,
            fgsz=(6, 6), ftsz=FTSZ, alpha=0.3, title="Scatter plot")

```



✓ Part 2. Plotting mission: The overtraining check



✓ 2.1 Understanding of the plot

Describe the plot and explain why this is called an "overtraining check" plot. Importance will be given to the clarity of your answer.

Explanation

The X-Axis shows the predicted probability given to each sample by the model, it serves different purposes according to the type of the data sample at hand. If it is a signal \rightarrow the model should predict a value close to 1.

If it is a Background \rightarrow the model should predict a value close to zero.

the filled histograms (blue and red) maps the predicted probabilities of the training data, and it shows good performance on the training set (because the majority of the prediction are both either close to one in the signal case, and close to zero in the background case)

the other part of the plot is the performance of the model on the test set, which is showed by the red and blue crosses. these dots or crosses do a similar behavior to the histograms, but for the test set. and as we can see, they reflect also a good performance on the test st from the model, and a huge correlation to the training set performance,

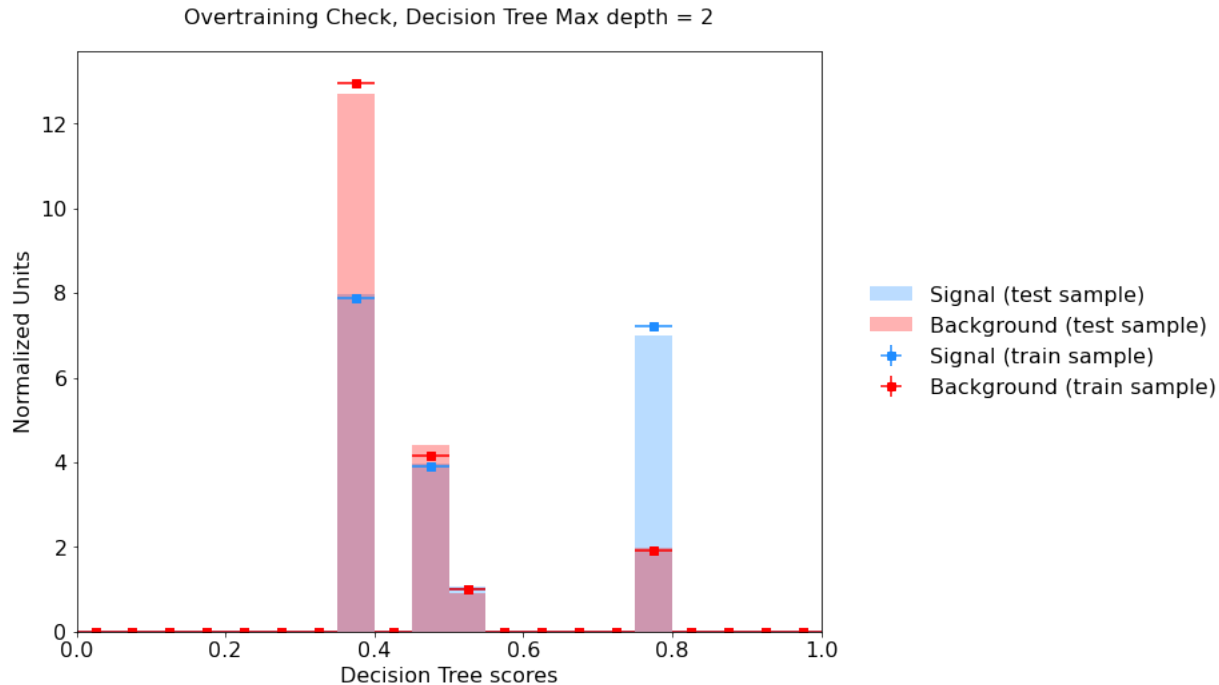
all in all, this pattern indicate a successfull training and generalization from the mode, as it reflects a good performance on unseen data(test set)

in case f over fitting, we should see a poor performance on the test set and a good performance on the training set.

✓ 2.2 Reproducing the plot

Write a function `plot_overtraining_check` that takes as arguments the classifier object, the and lists of the training and testing sets, the value of the positive class (e.g. for VBF it is 1) and a title. The function should split each dataset (train/test) into the real category (signal/background).

To test your plotting macro, use a Decision Tree classifier of maximum depth 2. You will obtain a plot like this:



```
def plot_overtraining_check(DT,Signal_train, Background_train,
    Signal_test, Background_test,
    xname=None, xlabel=None,
    yname=None, ylabel=None,
    figsize=(6, 6), fsize=16, alpha=0.3, title="Overtraining Check"):

    fig, ax = plt.subplots(figsize=figsize)

    # Annotate x-axis
    ax.set_xlabel(xlabel)

    # Annotate y-axis
    ax.set_ylabel(ylabel)

    # Training Set
    sig_train_predictions = DT.predict_proba(Signal_train[inputs])[:,1]
    bkg_train_predictions = DT.predict_proba(Background_train[inputs])[:,1]

    # Test set
    sig_test_predictions = DT.predict_proba(Signal_test[inputs])[:,1]
    bkg_test_predictions = DT.predict_proba(Background_test[inputs])[:,1]

    bins = np.linspace(0, 1.0, 21)

    # this gets the hieght of the training set histogram, (the position of the point )
    sig_train_hist, binEdges = np.histogram(sig_train_predictions, bins=bins, density=True)
    bkg_train_hist, binEdges = np.histogram(bkg_train_predictions, bins=bins, density=True)
```

```

bin_centers = 0.5 * (binEdges[1:]+binEdges[:-1])
# # Draw the test set histograms first ## THIS IS WHERE I NEEDED HELP FROM THE BRILLIANT TIM HITGE
ax.hist(sig_test_predictions, histtype='stepfilled', range= (0,1.0), bins = bins ,density=True, color = 'blue', alpha = 0.35, 1
ax.hist(bkg_test_predictions, histtype='stepfilled', range= (0,1.0), bins = bins ,density=True, color = 'red', alpha = 0.35, 1
# ax.errorbar(sig_test_predictions , sig_train_predictions,fmt = 'o', label = 'Signal - Train')
# ax.errorbar(bkg_train_predictions, histtype='stepfilled', range= (0,1.0), bins = 20 ,density=True, color = 'red', alpha = 0.
# Plot error bars (Mean predicted probability for each bin in training data)
ax.errorbar(bin_centers, sig_train_hist, yerr=np.sqrt(sig_train_hist/len(bkg_train_predictions)),
            fmt='o', color='blue', label='Signal - Train')
# both yerr = np.sqrt(y), and this version devided by the length can be used
ax.errorbar(bin_centers, bkg_train_hist, yerr=np.sqrt(bkg_train_hist / len(bkg_train_predictions)),
            fmt='o', color='red', label='Background - Train')
# https://stackoverflow.com/questions/11774822/matplotlib-histogram-with-errorbars

# Legend and plot:
ax.legend(fontsize=ftsize, bbox_to_anchor=(1.04, 0.5), loc="center left", frameon=False)
ax.set_title(title, pad=20)
plt.show()

```

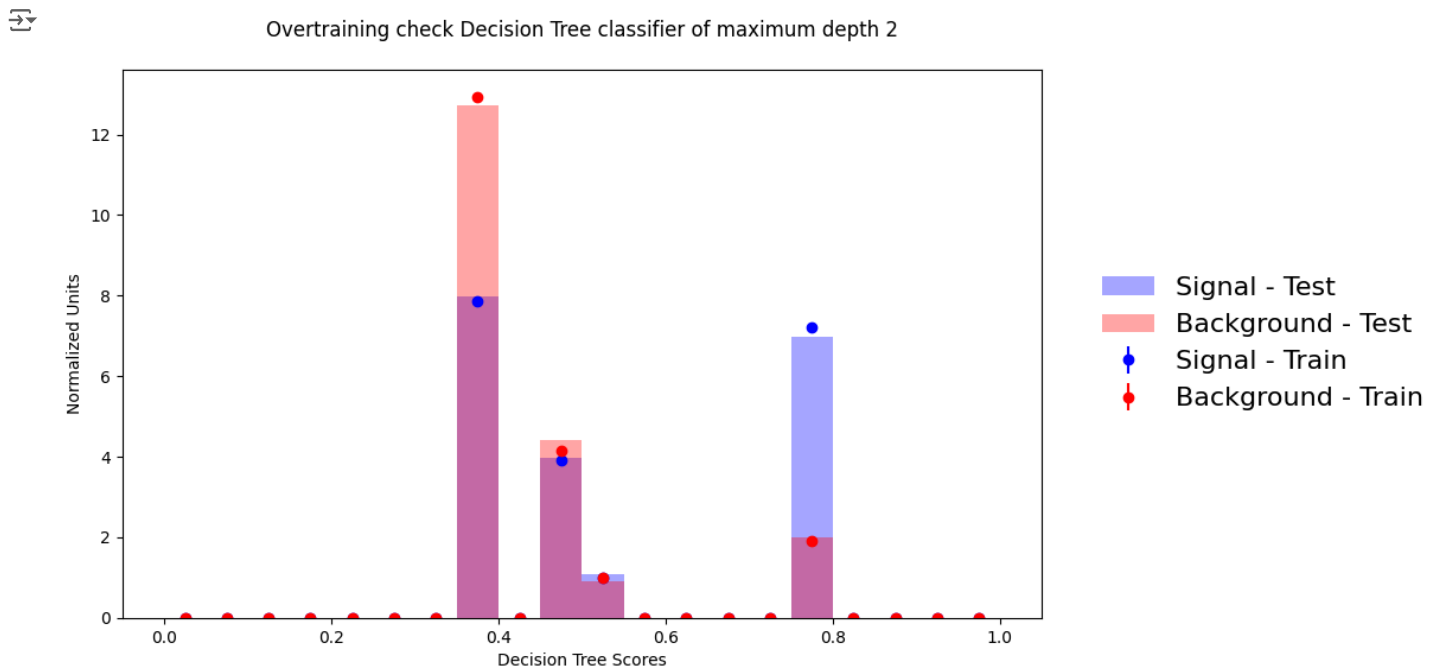
✓ Using it with Decision Tree classifier of maximum depth 2

```

from sklearn.tree import DecisionTreeClassifier
DT2 = DecisionTreeClassifier(max_depth=2)
DT2= DT2.fit(X_train, y_train)
# mostly done by Tim

plot_overtraining_check(DT2,Signal_train, Background_train,
                        Signal_test, Background_test,
                        xname=XNAME, xlabel='Decision Tree Scores',
                        yname=YNAME, ylabel='Normalized Units',
                        fgsize=(10, 6), ftsize=FTSIZE,title = "Overtraining check Decision Tree classifier of maximum depth 2",
                        alpha=0.3)

```



✓ 2.3 Using the plot

Create a random forest classifier with 100 estimators and leave other hyperparameters as default. Plot the overtraining check with this classifier. What are your observations? Is it classifying well on the training set? Is it under- or overtrain? Why?

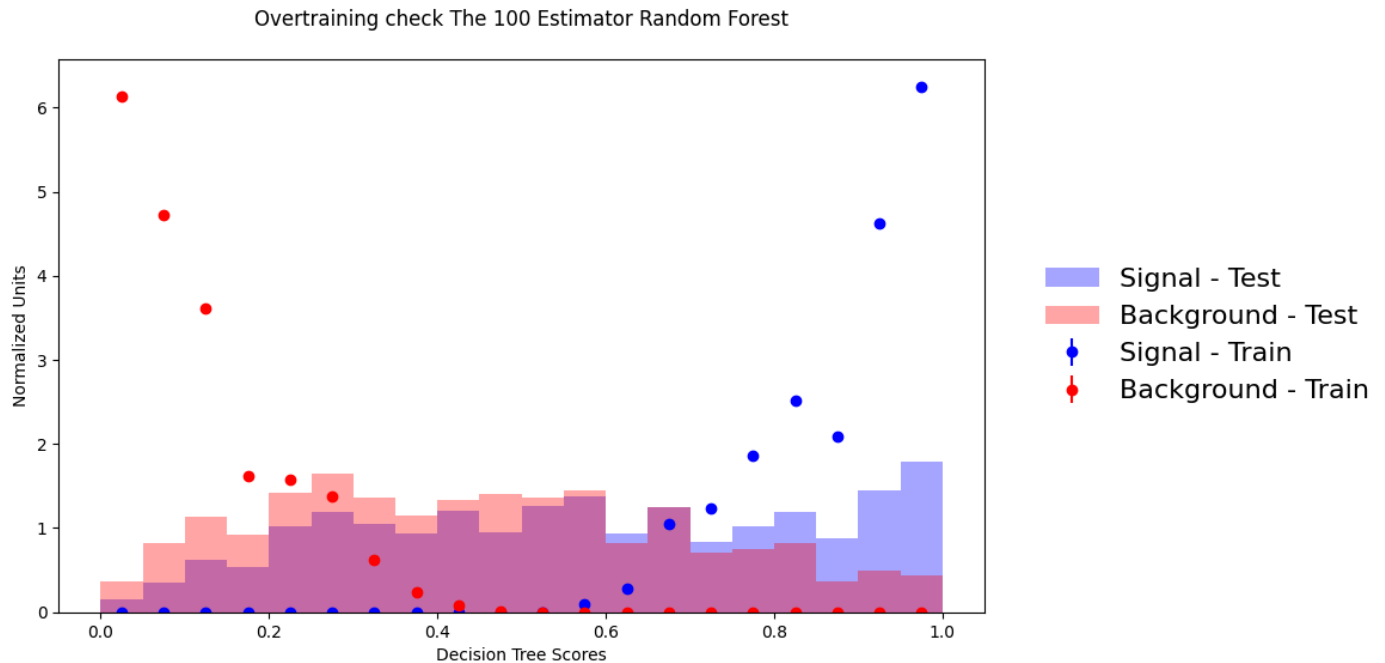
Create a second random forest classifier with this time the option `max_leaf_nodes=32`. What is improved? What is still problematic?

✓ The 100 Estimator Random Forest

```
from sklearn.ensemble import RandomForestClassifier
RF100 = RandomForestClassifier(n_estimators=100)
RF100= RF100.fit(X_train, y_train)
# mostly done by Tim
```

```
plot_overtraining_check(RF100,Signal_train, Background_train,
                        Signal_test, Background_test,
                        xname=XNAME, xlabel='Decision Tree Scores',
                        yname=YNAME, ylabel='Normalized Units',
                        fgsz=(10, 6), ftsz=FONTSIZE,title= "Overtraining check The 100 Estimator Random Forest",
                        alpha=0.3)
```

→ /usr/local/lib/python3.10/dist-packages/sklearn/base.py:1473: DataConversionWarning: A column-vector y was passed when a 1d return fit_method(estimator, *args, **kwargs)



My observation is that this classifier extremely overfits the training set because it performs well on the training set (the red and blue dots reflect accurate predictions)


But when it comes to the test set it performs very poorly as you can see, it nearly has a uniform distribution among the probabilities, reflecting no learning, and also a big contrast between the training and test performance

In short : Good performance on training set + poor performance on the test set -> overfitting


✓ The 32 Max leaf Random Forest

```
from sklearn.ensemble import RandomForestClassifier
RF32 = RandomForestClassifier(n_estimators=100, max_leaf_nodes=32)
RF32= RF32.fit(X_train, y_train)
# mostly done by Tim
```

```
plot_overtraining_check(RF32,Signal_train, Background_train,
                        Signal_test, Background_test,
                        xname=XNAME, xlabel='Decision Tree Scores',
                        yname=YNAME, ylabel='Normalized Units',
                        fgsz=(10, 6), ftsz=FONTSIZE, title = "Overtraining check The 32 Max leaf Random Forest" ,
                        alpha=0.3)
```

 /usr/local/lib/python3.10/dist-packages/sklearn/base.py:1473: DataConversionWarning: A column-vector y was passed when a 1d return fit_method(estimator, *args, **kwargs)

Overtraining check The 32 Max leaf Random Forest



Now thats a huge improvement!!!

the training performance is not perfect but good at the same time, a sign of no overfitting, the test set is also good with a clear pattern of separation between the negaive and postive sampse, but the biggest indication of no overfitting is the correlation between the train and test performances, indicating that the model is generalizing well to unseen data